

Atul Arun

352-219-9694 | atularun25@gmail.com | linkedin.com/in/atularun | github.com/Atul71

MS Computer Science student at University of Florida (GPA 4.0) with 4+ years deploying production ELT pipelines, Airflow orchestration, and Tableau/Power BI dashboards at enterprise scale. Proficient in Python, SQL, and Spark distributed infrastructure; experienced automating data quality checks and building analytics that drive operational decisions.

EDUCATION

University of Florida

Master of Science in Computer Science, GPA – 4.0/4.0

Gainesville, FL

Aug. 2025 – May 2027

- **Coursework:** Data Engineering, Introduction to Data Science, Distributed Systems, Database Systems

Model Engineering College, APJ Abdul Kalam Technological University

India

Bachelor of Technology in Computer Science Engineering, GPA – 9.52/10

Aug. 2017 – May 2021

TECHNICAL SKILLS

Languages: Python, SQL (advanced), PySpark, Scala (familiar), C++, JavaScript/TypeScript

Pipeline & Orchestration: Airflow, ELT/ETL design, Spark distributed infrastructure (PySpark), batch processing, data quality checks

Analytics & Dashboards: Tableau, Power BI, star schema modeling, KPI reporting, ad-hoc analysis

Cloud & DevOps: AWS (EKS, RDS, S3), Docker, Kubernetes, Jenkins, GitHub Actions, CI/CD, Git, Jira, Terraform

Databases: Oracle, SQL Server, NoSQL (MongoDB, DynamoDB), SQLite, ChromaDB

AI & ML: LLM API integration, RAG pipelines, MCP (Model Context Protocol), prompt engineering

Certifications: Microsoft Certified Fabric Data Engineer Associate (DP-700) | Power BI Data Analyst Associate

EXPERIENCE

UF Innovate

Patent & Technology Research Intern (AI, Software)

Apr. 2026 – Present

Gainesville, FL

- Built a Python pipeline to run completed commercialization reports through an LLM API and automated **Power BI dashboards** to track portfolio KPIs, cutting analyst turnaround by **50%** and enabling self-serve insights.

Bosch

Senior Data Engineer

Jul. 2024 – Jul. 2025

India

- Built an **ELT pipeline** integrating enterprise project data with an LLM API for automated capacity analysis; identified a **20% resource surplus** across **25,000+ employees**, directly informing staffing decisions for **500+ managers** globally.
- Designed and deployed a Python-based schema provisioning engine executing **4,300+ SQL queries in 48 hours** against Oracle DB, eliminating a 30-day manual bottleneck (**93% reduction** in deployment time).
- Resolved a data ingestion bottleneck by redesigning the logging write path and automating SQL retention jobs, cutting daily ingest **60%** (4 GB → 1.5 GB) and eliminating storage throttling.

Data Engineer

Sep. 2021 – Jun. 2024

- Deployed **Airflow-scheduled ELT pipelines** consolidating data from **5+ source systems** (SAP, Oracle, REST APIs) into a star-schema model using **Spark distributed infrastructure**; automated data quality checks reduced pipeline failures by **90%**.
- Built **Tableau and Power BI dashboards** for **400+ users** across **50+ plants**; consolidated 40 fragmented reports into a governed star-schema model, improving reporting efficiency by **30%**.
- Standardized **Git and Jira** workflows across 4 domain teams; migrated 40 Power BI reports to **Azure DevOps** enabling full rollback on pipeline releases.

Data & Analytics Intern

Mar. 2021 – Jun. 2021

- Built a job monitoring dashboard with PL/SQL-driven alerts for **550 QM jobs**, reducing stakeholder-reported failures by **90%** by surfacing anomalies before escalation.
- Applied statistical trend and variance analysis across **50+ plants** to identify cost drivers, improving forecasting accuracy by **20%**.

PROJECTS

Self-Healing Civic Data Pipeline | Python, MCP, Airflow, RAG, ChromaDB, SQLite

[Link](#)

- Built an ELT pipeline with **automated data quality repair**: LLM-supervised MCP layer fixes corrupted NYC311 records via geospatial inference and RAG, raising F1 from **0.11** to **0.63** (**5.7×** over deterministic baseline) on a 1,000-row benchmark.

MCP-Orchestrated ELT Pipeline | Python, MCP, PySpark, SQLite, GitHub Actions

[Link](#)

- Built a modular ELT pipeline for Gainesville crime data using three MCP servers (extract, transform, load); **Pytest data quality suite** and **GitHub Actions CI/CD** ensure schema-validated, deployment-ready output on every commit.

Distributed Reddit Engine | Gleam, Erlang/OTP, Mist, REST

[Link](#)

- Architected a fault-tolerant distributed backend on the Erlang/OTP actor model; maintained **atomic state consistency** across nodes under failure with graceful degradation.